ASSESMENT 2

ASSESMENT DATE	26-09-2022
STUDENT NAME	P.Keerthi
STUDENT ROLL NUMBER	713119104005
MAXIMUM MARKS	2 Marks

1. **Download the dataset**
2. **Load the dataset Solution:**

   import pandas as pd
   import numpy as np
   df=pd.read_csv("C:\\Users\\PC\\Desktop\\Churn_Modelling.csv")

df

output:

df.head()

3.perform following operations

   ➢ **univariate analysis**

```
import matplotlib.pyplot as plt
import seaborn as sns
plt.scatter(df.index,df['CreditScore'])
plt.show()
```

output:
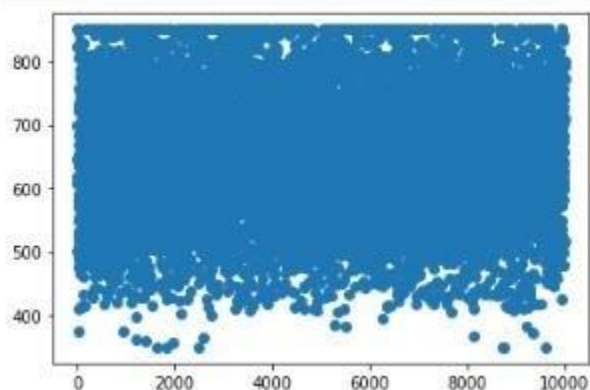
plt.hist(df['Tenure'])

```
Out[12]: (array([ 413., 1035., 1048., 1009.,  989., 1012.,  967., 1028., 1025.,
                 1474.]),
         array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.]),
         <a list of 10 Patch objects>)
```


sns.distplot(df['NumOfProducts'])

plt.boxplot(df['RowNumber'])

plt.show()

sns.boxplot(df['CustomerId'])

df['HasCrCard'].value_counts().plot.bar()

sns.countplot(df['Geography'])

> **Bivariate analysis**

plt.scatter(df.Gender, df.Age)

plt.title('employee genger vs. employee

age') plt.xlabel('employee details')

plt.ylabel('range')

employee genger vs. employee age
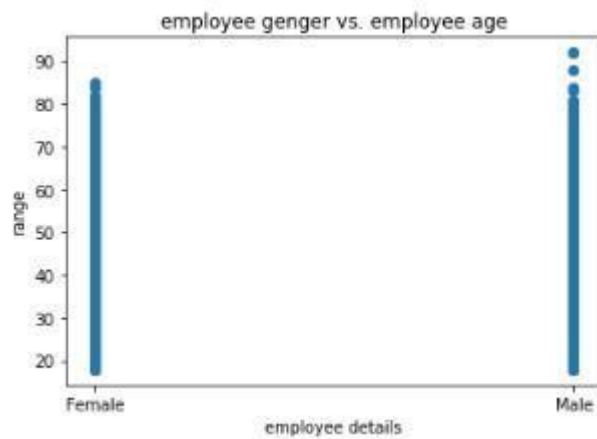
plt.scatter(df.NumOfProducts, df.Balance)

plt.title('numof products vs balance')

plt.xlabel('graph for numof products')

plt.ylabel('graph for balance')


```
Out[20]: Text(0, 0.5, 'graph for balance')
```

numof products vs balance
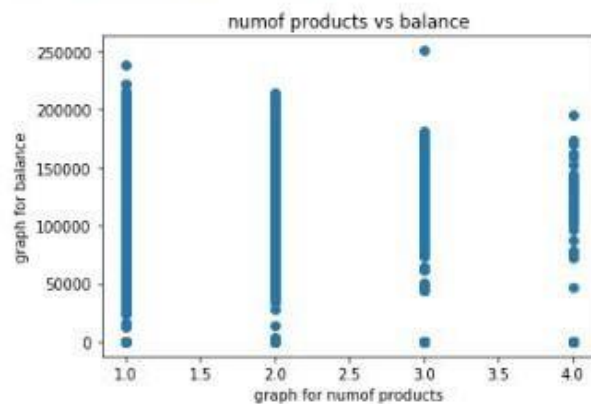
> **Multivariate analysis**

seaborn.pairplot(df)

plt.show()
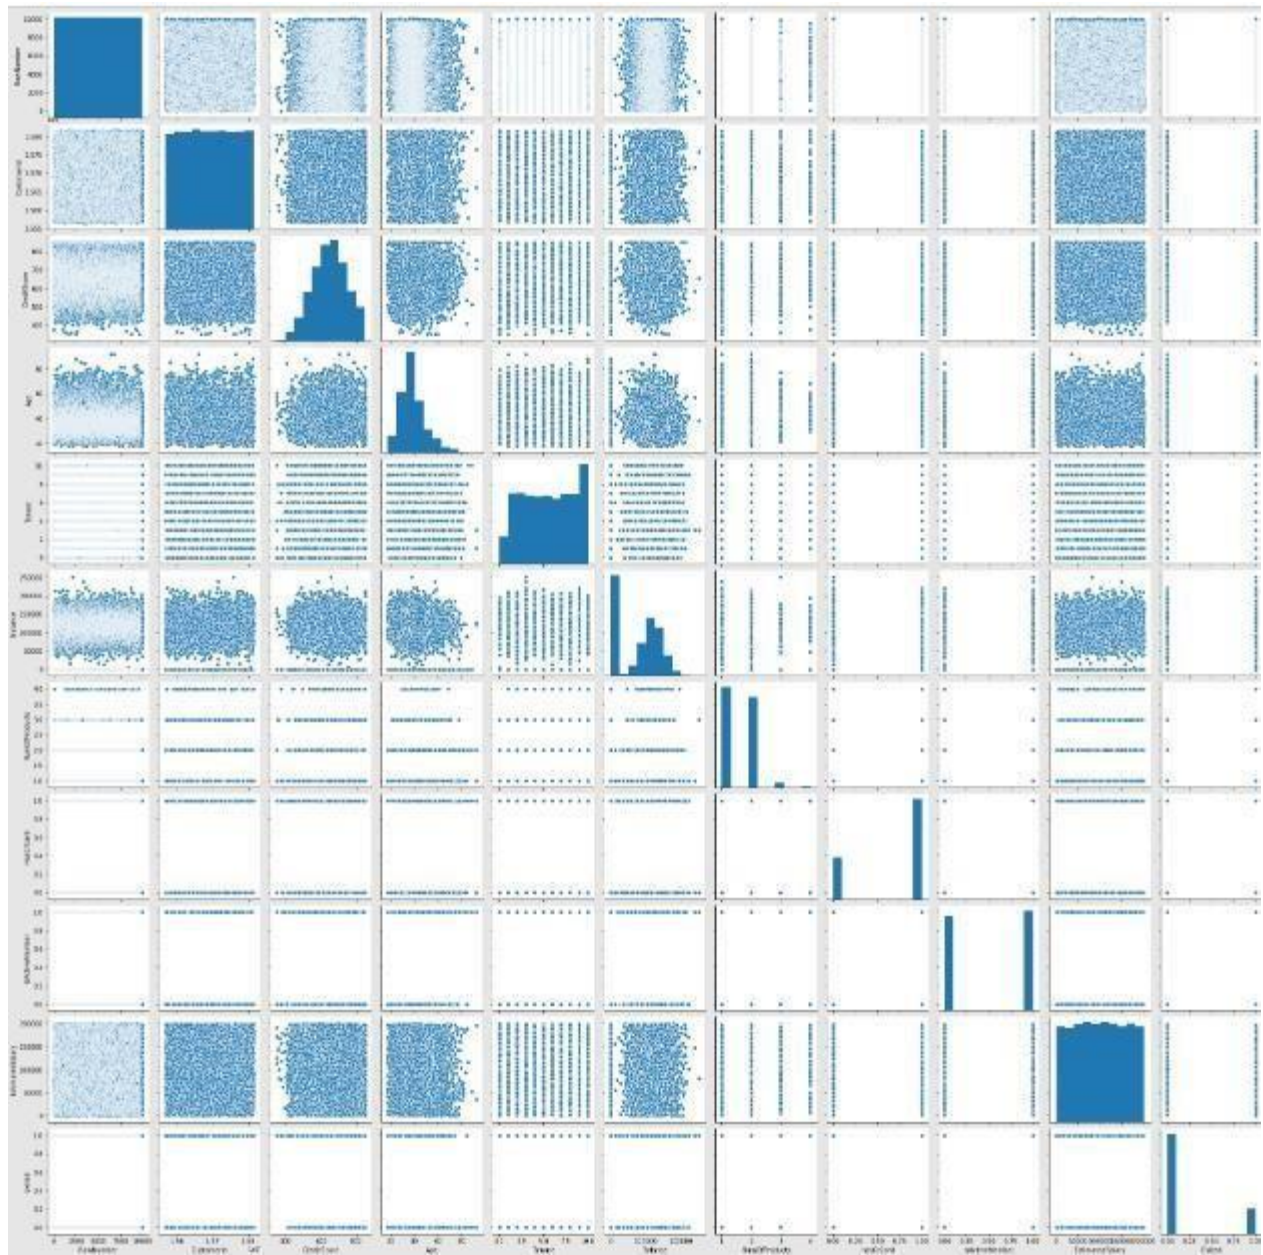
4.describtive function

df.describe()

Out[6]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000	10
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000	
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000	
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000	
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500	
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000	

5.handle the missing data df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   RowNumber        10000 non-null  int64
 1   CustomerId       10000 non-null  int64
 2   Surname          10000 non-null  object
 3   CreditScore      10000 non-null  int64
 4   Geography        10000 non-null  object
 5   Gender           10000 non-null  object
 6   Age              10000 non-null  int64
 7   Tenure           10000 non-null  int64
 8   Balance          10000 non-null  float64
 9   NumOfProducts    10000 non-null  int64
 10  HasCrCard        10000 non-null  int64
 11  IsActiveMember   10000 non-null  int64
 12  EstimatedSalary  10000 non-null  float64
 13  Exited           10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

df.isnull()

Out[14]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSal
0	False	False	False	False	False	False	False	False	False	False	False	False	Fa
1	False	False	False	False	False	False	False	False	False	False	False	False	Fa
2	False	False	False	False	False	False	False	False	False	False	False	False	Fa
3	False	False	False	False	False	False	False	False	False	False	False	False	Fa
4	False	False	False	False	False	False	False	False	False	False	False	False	Fa
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9996	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9997	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9998	False	False	False	False	False	False	False	False	False	False	False	False	Fa
9999	False	False	False	False	False	False	False	False	False	False	False	False	Fa

10000 rows × 14 columns

df.notnull()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSala
0	True	True	True	True	True	True	True	True	True	True	True	True	Tru
1	True	True	True	True	True	True	True	True	True	True	True	True	Tru
2	True	True	True	True	True	True	True	True	True	True	True	True	Tru
3	True	True	True	True	True	True	True	True	True	True	True	True	Tru
4	True	True	True	True	True	True	True	True	True	True	True	True	Tru
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	True	True	True	True	True	True	True	True	True	True	True	True	Tru
9996	True	True	True	True	True	True	True	True	True	True	True	True	Tru
9997	True	True	True	True	True	True	True	True	True	True	True	True	Tru
9998	True	True	True	True	True	True	True	True	True	True	True	True	Tru
9999	True	True	True	True	True	True	True	True	True	True	True	True	Tru

10000 rows × 14 columns

df.fillna(0)

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSa
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows × 14 columns

df["Gender"].fillna("No Gender", inplace = True) df

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSa
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	10134
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	11254
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	11393
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	9382
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	7908
...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijiaku	771	France	Male	39	5	0.00	2	1	0	9627
9996	9997	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	10169
9997	9998	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	4208
9998	9999	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	9288
9999	10000	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	3819

10000 rows × 14 columns

df.drop("RowNumber",axis=1,inplace=True)

df

	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10
...	...	...	...	...	...	...	...	...	...	...	...	...
9995	15606229	Obijiaku	771	France	Male	39	5	0.00	2	1	0	96270.64
9996	15569892	Johnstone	516	France	Male	35	10	57369.61	1	1	1	101699.77
9997	15584532	Liu	709	France	Female	36	7	0.00	1	0	1	42085.58
9998	15682355	Sabbatini	772	Germany	Male	42	3	75075.31	2	1	0	92888.52
9999	15628319	Walker	792	France	Female	28	4	130142.79	1	1	0	38190.78

10000 rows × 12 columns

print(df.isnull().sum())

```
CustomerId         0
Surname            0
CreditScore        0
Geography          0
Gender             0
Age                0
Tenure             0
Balance            0
NumOfProducts      0
HasCrCard          0
IsActiveMember     0
EstimatedSalary    0
dtype: int64
```

updated_df = df.dropna(axis=1) updated_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   CustomerId       10000 non-null  int64
 1   Surname          10000 non-null  object
 2   CreditScore      10000 non-null  int64
 3   Geography        10000 non-null  object
 4   Gender           10000 non-null  object
 5   Age              10000 non-null  int64
 6   Tenure           10000 non-null  int64
 7   Balance          10000 non-null  float64
 8   NumOfProducts    10000 non-null  int64
 9   HasCrCard        10000 non-null  int64
 10  IsActiveMember   10000 non-null  int64
 11  EstimatedSalary  10000 non-null  float64
dtypes: float64(2), int64(7), object(3)
memory usage: 937.6+ KB
```

6. Finding outliers and replace

Q1 = df.quantile(0.25)

Q3 = df.quantile(0.75) IQR

= Q3 - Q1

print(IQR)

```
RowNumber            4999.5000
CustomerId         124705.5000
CreditScore           134.0000
Age                    12.0000
Tenure                  4.0000
Balance            127644.2400
NumOfProducts           1.0000
HasCrCard               1.0000
IsActiveMember          1.0000
EstimatedSalary     98386.1375
Exited                  0.0000
dtype: float64
```

print(df < (Q1 - 1.5 * IQR))

(df > (Q3 + 1.5 * IQR))

```
        Age   Balance  CreditScore  CustomerId  EstimatedSalary  Exited  \
0     False     False        False       False            False   False
1     False     False        False       False            False   False
2     False     False        False       False            False   False
3     False     False        False       False            False   False
4     False     False        False       False            False   False
...     ...       ...          ...         ...              ...     ...
9995  False     False        False       False            False   False
9996  False     False        False       False            False   False
9997  False     False        False       False            False   False
9998  False     False        False       False            False   False
9999  False     False        False       False            False   False

      Gender  Geography  HasCrCard  IsActiveMember  NumOfProducts  RowNumber  \
0      False      False      False           False          False      False
1      False      False      False           False          False      False
2      False      False      False           False          False      False
3      False      False      False           False          False      False
4      False      False      False           False          False      False
...      ...        ...        ...             ...            ...        ...
9995   False      False      False           False          False      False
9996   False      False      False           False          False      False
9997   False      False      False           False          False      False
9998   False      False      False           False          False      False
9999   False      False      False           False          False      False

      Surname  Tenure
0       False   False
1       False   False
2       False   False
3       False   False
4       False   False
...       ...     ...
9995    False   False
9996    False   False
9997    False   False
9998    False   False
9999    False   False
```

plt.boxplot(df["NumOfProducts"])

plt.show()

np.where(df.Age>42,42, df.Age)

```
Out[16]: array([42, 41, 42, ..., 36, 42, 28], dtype=int64)
```

print(df['Age'].skew())

```
1.0113202630234552
```

print(df['Age'].quantile(0.25))

print(df['Age'].quantile(0.75))

df['Age'] = np.where(df['Age'] > 39, 41, df['Age'])

df.describe()

```
32.0
41.0
```

Out[22]:		RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
	count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	10000.000000	10000.000000
	mean	5000.50000	1.569094e+07	650.528800	35.788600	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881
	std	2886.89568	7.193619e+04	96.653299	5.659409	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818
	min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000	0.000000	11.580000
	25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000	0.000000	51002.110000
	50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000	1.000000	100193.915000
	75%	7500.25000	1.575323e+07	718.000000	41.000000	7.000000	127644.240000	2.000000	1.00000	1.000000	149388.247500
	max	10000.00000	1.581569e+07	850.000000	41.000000	10.000000	250898.090000	4.000000	1.00000	1.000000	199992.480000

7.categorical column

df["CustomerId"].value_counts()

```
Out[23]: 15812607    1
         15741078    1
         15635776    1
         15740223    1
         15738174    1
                    ..
         15743714    1
         15639265    1
         15641312    1
         15684319    1
         15695872    1
         Name: CustomerId, Length: 10000, dtype: int64
```

df.dtypes

```
Out[27]: RowNumber           int64
         CustomerId          int64
         Surname          category
         CreditScore         int64
         Geography          object
         Gender             object
         Age                 int64
         Tenure              int64
         Balance           float64
         NumOfProducts       int64
         HasCrCard           int64
         IsActiveMember      int64
         EstimatedSalary   float64
         Exited              int64
         dtype: object
```

df["Age"].value_counts().sort_index()

```
Out[32]: 18      22
         19      27
         20      40
         21      53
         22      84
         23      99
         24     132
         25     154
         26     200
         27     209
         28     273
         29     348
         30     327
         31     404
         32     418
         33     442
         34     447
         35     474
         36     456
         37     478
         38     477
         39     423
         41    4013
         Name: Age, dtype: int64
```

df_categorical = df[categorical_columns]

df_categorical.head()

Out[35]:

	Geography	Gender
0	France	Female
1	Spain	Female
2	France	Female
3	France	Female
4	Spain	Female

pd.get_dummies(df, columns=["Age"]).head()

Out[36]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Tenure	Balance	NumOfProducts	HasCrCard	...	Age_31	Age_32	Age_33	Age_34
0	1	15634602	Hargrave	619	France	Female	2	0.00	1	1	...	0	0	0	0
1	2	15647311	Hill	608	Spain	Female	1	83807.86	1	0	...	0	0	0	0
2	3	15619304	Onio	502	France	Female	8	159660.80	3	1	...	0	0	0	0
3	4	15701354	Boni	699	France	Female	1	0.00	2	0	...	0	0	0	0
4	5	15737888	Mitchell	850	Spain	Female	2	125510.82	1	1	...	0	0	0	0

5 rows × 36 columns

8.split the data into dependent and independent variables print(df.size)

```
140000
```

X = df.iloc[:, :-1].values

print(X)

```
[[1 15634602 'Hargrave' ... 1 1 101348.88]
 [2 15647311 'Hill' ... 0 1 112542.58]
 [3 15619304 'Onio' ... 1 0 113931.57]
 ...
 [9998 15584532 'Liu' ... 0 1 42085.58]
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]
 [10000 15628319 'Walker' ... 1 0 38190.78]]
```

Y = df.iloc[:, -1].values

print(Y)

```
[1 0 1 ... 1 1 0]
```

9.minmaxscaler

from sklearn.preprocessing import MinMaxScaler

df

scaler = MinMaxScaler()

print(scaler.fit(df))

```
MinMaxScaler(copy=True, feature_range=(0, 1))
```

10. train –split data import pandas as pd from sklearn.linear_model import LinearRegression from sklearn.model_selection import train_test_split

df=pd.read_csv("C:\\Users\\PC\\Desktop\\Churn_Modelling.csv")

df.head()

Out[77]:	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

y= df.Tenure

y.head()

```
Out[78]: 0    2
         1    1
         2    8
         3    1
         4    2
         Name: Tenure, dtype: int64
```

x=df.drop('Tenure',axis=1)

x.head()

Out[80]:	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	France	Female	42	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	Spain	Female	43	125510.82	1	1	1	79084.10	0

x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)

x_train.shape

```
Out[82]: (8000, 13)
```

y_train.shape

```
Out[83]: (8000,)
```

x_test.shape

```
Out[84]: (2000, 13)
```

y_test.shape

```
Out[85]: (2000,)
```